

Prediction of Diabetes in Women

By Aparajita Mukherjee

Table of Content

| | | |
|-------|--|----|
| 1. | Project Objective | 3 |
| 2. | Details of the Project | 3 |
| 3. | Required Packages..... | 4 |
| 4. | Basic EDA(Exploratory Data Analysis) | 4 |
| 4.1 | Missing Values | 5 |
| 4.2 | Outliers | 6 |
| 4.3 | Correlation | 6 |
| 4.4 | Visualization(Independent vs Dependent) | 8 |
| 4.4.1 | Scatter Plot | 8 |
| 4.4.2 | Boxplot | 10 |
| 5. | Model Building..... | 13 |
| 5.1 | Split the Data | 13 |
| 5.2 | Logistic Regression..... | 13 |
| 5.2.1 | Model Building | 13 |
| 5.2.2 | Interpretation | 14 |
| 5.2.3 | Multicollinearity Check..... | 14 |
| 5.2.4 | Improving the Model(AIC) | 15 |
| 5.2.5 | Tuning the Model(to improve Sensitivity) | 18 |
| 5.3 | K-NN(K-Nearest Neighbour) | 19 |
| 5.3.1 | Normalizing continuous variables | 19 |
| 5.3.2 | Model Building | 19 |
| 5.4 | Naïve Bayes | 21 |
| 5.4.1 | Model Building | 21 |
| 6. | Model Comparison | 22 |
| 6.1 | Logistic Regression..... | 22 |
| 6.1.1 | Train Dataset | 22 |
| 6.1.2 | Test Data..... | 22 |
| 6.2 | K Nearest Neighbour..... | 23 |
| 6.2.1 | Train Dataset | 23 |
| 6.2.2 | Test Data..... | 24 |
| 6.3 | Naïve Bayes | 24 |
| 6.3.1 | Train Dataset | 24 |
| 6.3.2 | Test Data..... | 25 |
| 7. | Cross Validation & Insights..... | 26 |
| 7.1 | Cross Validation | 26 |
| 7.2 | Insights | 27 |
| 8. | Source of Data | 27 |

1. Project Objective

The Objective of this project is to –

1. Plot graphs between Dependent and Independent Variables.
2. Split the data into in-sample(train) and out-of-sample(test) – 70%-30%.
3. Build models using – Naïve Bayes, K-NN, Logistic Regression on train dataset.
4. Validate the model on the test dataset(to check for overfitting).
5. Calculate accuracy using Confusion Matrix.
6. Compare the results and draw inferences.

2. Details of the Project

A total of 9 variables were used for segregating the cases where there were occurrence of diabetes as compared to cases where there were none.

| Variable Name | Data Type | Variable Description |
|---------------|-----------|--|
| NoPreg | integer | Number of time pregnant |
| PlaGluConc | integer | Plasma glucose concentration a 2 hours in an oral glucose tolerance test |
| DiastolicBP | integer | Diastolic blood pressure (mm Hg) |
| TSkinThick | integer | Triceps skin fold thickness (mm) |
| Test | integer | 2-Hour serum insulin (μ U/ml) |
| BMI | numeric | Body mass index ($\text{weight in kg}/(\text{height in m})^2$) |
| DiabPediFunc | numeric | Diabetes pedigree function |
| Age | integer | Age (years) |
| Class | integer | Class variable (0 or 1) |

The datatype of the Class variable(target variable) requires to be changed to factor before using it as input to the logistic regression model.

Also, the dataset needs to be normalized before using it to build the KNN(K-Nearest Neighbour model and Naïve Bayes.

3. Required Packages

| Library | Description |
|-----------------------|--|
| library(ggplot2) | Create Data Visualisations |
| library(car) | Companion to Applied Regression |
| library(caret) | Classification and Regression Tree |
| library(class) | Functions for Classification |
| library(devtools) | Tools to make developing R packages easier |
| library(e1071) | Misc Functions of Statistics(Naïve Bayes) |
| library(lmtest) | Testing Linear Regression Models |
| library(Hmisc) | Get the detailed summary of the dataset |
| library(ROCR) | Visualizing the Performance of Scoring Classifiers |
| library(plyr) | Tools for Splitting, Applying & Combining Data |
| library(pROC) | Display & Analyse ROC curve |
| library(psych) | Procedures for Psychometric and Personality Research |
| library(dplyr) | A Grammar for Data Manipulation |
| library(corrplot) | Visualization of Correlation Matrix |
| library(caTools) | Splitting the Dataset |
| library(DataExplorer) | Automate Data Exploration and Treatment |

4. Basic EDA(Exploratory Data Analysis)

| | n | nmiss | outlier_flag | mean | stdev | min | q1.1% | q5.5% | q95.95% | q99.99% | max | UC | LC |
|--------------|-----|-------|--------------|-------------|-------------|--------|----------|----------|-----------|-----------|--------|------------|--------------|
| NoPreg | 768 | 0 | 1 | 3.8450521 | 3.3695781 | 0.000 | 0.00000 | 0.00000 | 10.00000 | 13.00000 | 17.00 | 13.953786 | -6.2636821 |
| PlaGluConc | 768 | 0 | 1 | 120.8945312 | 31.9726182 | 0.000 | 57.00000 | 79.00000 | 181.00000 | 196.00000 | 199.00 | 216.812386 | 24.9766767 |
| DiastolicBP | 768 | 0 | 1 | 69.1054688 | 19.3558072 | 0.000 | 0.00000 | 38.70000 | 90.00000 | 106.00000 | 122.00 | 127.172890 | 11.0380472 |
| TSkinThick | 768 | 0 | 1 | 20.5364583 | 15.9522176 | 0.000 | 0.00000 | 0.00000 | 44.00000 | 51.33000 | 99.00 | 68.393111 | -27.3201944 |
| Test | 768 | 0 | 1 | 79.7994792 | 115.2440024 | 0.000 | 0.00000 | 0.00000 | 293.00000 | 519.90000 | 846.00 | 425.531486 | -265.9325279 |
| BMI | 768 | 0 | 1 | 31.9925781 | 7.8841603 | 0.000 | 0.00000 | 21.80000 | 44.39500 | 50.75900 | 67.10 | 55.645059 | 8.3400972 |
| DiabPediFunc | 768 | 0 | 1 | 0.4718763 | 0.3313286 | 0.078 | 0.09468 | 0.14035 | 1.13285 | 1.69833 | 2.42 | 1.465862 | -0.5221095 |
| Age | 768 | 0 | 1 | 33.2408854 | 11.7602315 | 21.000 | 21.00000 | 21.00000 | 58.00000 | 67.00000 | 81.00 | 68.521580 | -2.0398092 |
| Class | 768 | 0 | 0 | 0.3489583 | 0.4769514 | 0.000 | 0.00000 | 0.00000 | 1.00000 | 1.00000 | 1.00 | 1.779812 | -1.0818958 |

As is clear from the above table that –

- There are 9 variables having 768 observations.
- There are no missing values.
- All the 8 variables except for the target variable, have outliers.

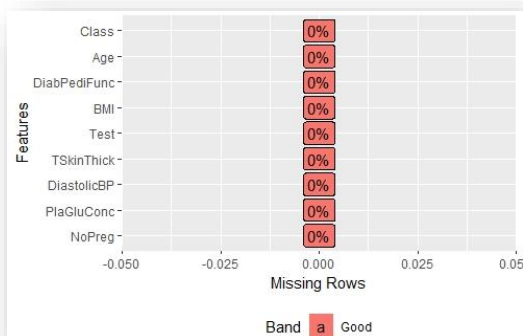
```
> str(data)
'data.frame': 768 obs. of 9 variables:
 $ NoPreg      : int  6 1 8 1 0 5 3 10 2 8 ...
 $ PlaGluConc  : int 148 85 183 89 137 116 78 115 197 125 ...
 $ DiastolicBP : int  72 66 64 66 40 74 50 0 70 96 ...
 $ TSkinThick  : int  35 29 0 23 35 0 32 0 45 0 ...
 $ Test        : int  0 0 0 94 168 0 88 0 543 0 ...
 $ BMI         : num 33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
 $ DiabPediFunc: num 0.627 0.351 0.672 0.167 2.288 ...
 $ Age         : int  50 31 32 21 33 30 26 29 53 54 ...
 $ Class       : int  1 0 1 0 1 0 1 0 1 1 ...
```

```
> summary(data)
      NoPreg      PlaGluConc      DiastolicBP      TSkinThick      Test
Min.   : 0.000   Min.   : 0.0   Min.   : 0.00   Min.   : 0.00   Min.   : 0.0
1st Qu.: 1.000   1st Qu.: 99.0   1st Qu.: 62.00  1st Qu.: 0.00  1st Qu.: 0.0
Median : 3.000   Median :117.0   Median : 72.00  Median :23.00  Median : 30.5
Mean   : 3.845   Mean   :120.9   Mean   : 69.11  Mean   :20.54  Mean   : 79.8
3rd Qu.: 6.000   3rd Qu.:140.2   3rd Qu.: 80.00  3rd Qu.:32.00  3rd Qu.:127.2
Max.   :17.000   Max.   :199.0   Max.   :122.00  Max.   :99.00  Max.   :846.0

      BMI      DiabPediFunc      Age      Class
Min.   : 0.00   Min.   :0.0780   Min.   :21.00   Min.   :0.000
1st Qu.:27.30   1st Qu.:0.2437   1st Qu.:24.00   1st Qu.:0.000
Median :32.00   Median :0.3725   Median :29.00   Median :0.000
Mean   :31.99   Mean   :0.4719   Mean   :33.24   Mean   :0.349
3rd Qu.:36.60   3rd Qu.:0.6262   3rd Qu.:41.00   3rd Qu.:1.000
Max.   :67.10   Max.   :2.4200   Max.   :81.00   Max.   :1.000
```

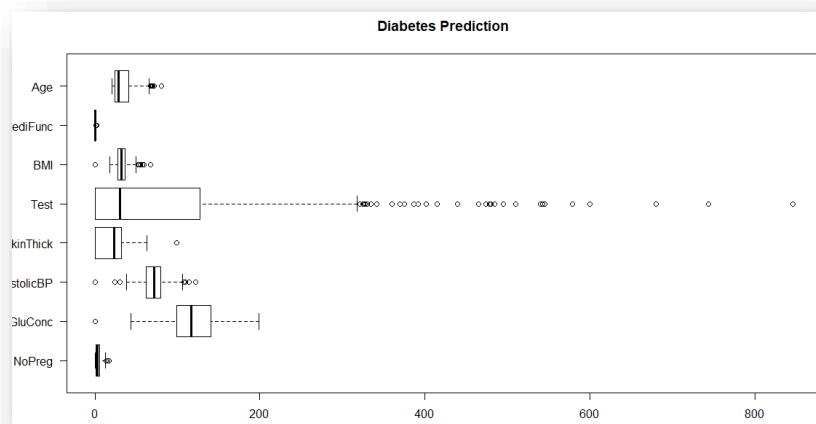
- There are 9 variables having 768 observations.
- There are evidences of outliers in the dataset which will be clear after plotting the same through a boxplot.

4.1 Missing Values




Missing values in the dataset can lower the accuracy rate of the model. It is therefore necessary to either remove the missing values (if the percentage of missing values is very small) or replace them with the Mean, Median or Mode of the particular variable. From the above plot it is evident that there are no missing values in the dataset and it is good to start working with.

4.2 Outliers



The plot indicates that there are outliers in all the independent variables of the dataset. But the algorithms that will be used will not be significantly affected by the presence of outliers. In case, necessary, we will cap the outlier values to the 99th quantile of that particular variable.


| Diabetes | Count | Percentage |
|----------|-------|------------|
| 0 | 500 | 65% |
| 1 | 268 | 35% |
| Total | 768 | 100% |



The number of women detected to be suffering from diabetes is very less as compared to women who were not suffering from it.

4.3 Correlation


It is evident that there is not much correlation between the variables. This wouldn't affect the accuracy of the model.



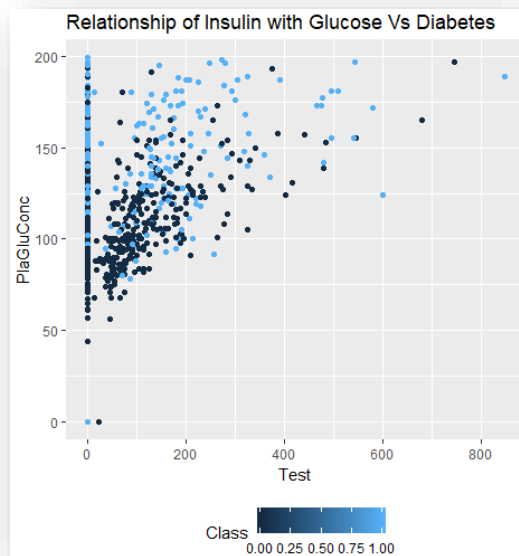
- There is no obvious relationship between age and onset of diabetes.
- There is no obvious relationship between pedi function and onset of diabetes.
- This may suggest that diabetes is not hereditary, or that the Diabetes Pedigree Function needs work.
- Variables like PleGluConc, NoPreg, BMI, Age, DiabPediFunc and Test have a positive relation with the target variable(Class). This means whenever the value for these variables increases, the minority class also increases.
- Most of the variables are positively related.
-
- Larger values of plas combined with larger values for age, pedi, mass, insu, skin, pres, and preg tends to show greater likelihood of testing positive for diabetes.

4.4 Visualization(Independent vs Dependent)


4.4.1 Scatter Plot



No clear boundary can be drawn that separates Non-diabetic and Diabetic women based on Number of Pregnancies vs Age




Non-diabetic women seemed to have lower levels of Insulin and Glucose as opposed to Diabetic women who recorded low to high levels of Insulin and high levels of Glucose




There is no significant distinction based on Insulin level and Pregnancies, but diabetic women seem to have a slightly increased insulin and the Pregnancies.



Non-diabetic women seemed to have lower BP levels and Age as opposed to Diabetic women who recorded low to high levels of BP and at a later Age.




Women who have Diabetes can be differentiated from those who don't have based on BMI and BP values




Women with low values of BMI and Skin Thickness did not have Diabetes


4.4.2 Boxplot




Higher the number of times a women was pregnant, she will have more probability to be diagnosed with Diabetes. This has a lot of data to the right side of the median as the 1st quartile is more than the 3rd quartile.




Higher the number of times a wo




The boxplot looks almost symmetric and relatively uniform. But a woman with marginally higher Diastolic Blood Pressure will be suffering from diabetes.




The value of the 1st quartile for the variable – TskinThick is “0” and therefore, the boxplot starts at the same point. It also indicates that women with marginally higher Triceps Skin Fold Thickness tend to suffer from diabetes.




As was seen in the summary, the minimum and the value for the first quartile for the variable – Test starts from “0”, especially for the women who were detected with the diabetes.



Shows that all the women who had Diabetes had a BMI greater than 25, which is above the normal levels. On the other hand, women who did not have Diabetes had a BMI ranging from 18 to 60.



Women with marginally higher amount of Diabetes Pedigree Function are more susceptible to having diabetes.



This boxplot indicates that higher age is a major factor contributing to the cause of diabetes in women.

5. Model Building

Algorithms used –

- Logistic Regression
- K-Nearest Neighbour
- Naïve Bayes

These algorithms are relevant because they perform classification on a dataset, deal appropriately with missing or erroneous data.

5.1 Split the Data

Before using the dataset for building and evaluating the models, it is advisable to split the data. The default ratio to split is usually a 70-30 split.

```
> dim(train_data)
[1] 534  9
> dim(test_data)
[1] 234  9
```

We want to ensure that the proportion of overall class variable is as relative as it was in the original dataset and the proportion gets retained in the train and test dataset.

```
> prop.table(table(data$Class))
      0      1
0.6510417 0.3489583
> prop.table(table(train_data$Class))
      0      1
0.6610487 0.3389513
> prop.table(table(test_data$Class))
      0      1
0.6282051 0.3717949
```

5.2 Logistic Regression

5.2.1 Model Building

A full model was built with Class as the response variable with the rest of the 8 predictor variables. Step-wise variable selection method was used to identify the most important variables. The final model chosen with AIC as the criterion for selection generated a logistic regression model with the lowest AIC value of 532.94 as below.

```

> logit_model1 = glm(Class ~ ., data = train_data,
+ family = binomial)
> summary(logit_model1)

Call:
glm(formula = Class ~ ., family = binomial, data = train_data)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.5420  -0.7458  -0.4474   0.7747   2.8993

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -7.547214   0.817328  -9.234 < 2e-16 ***
NoPreg       0.143319   0.038256   3.746 0.000179 ***
PlasGluConc  0.032700   0.004342   7.531 5.04e-14 ***
DiastolicBP -0.010214   0.006172  -1.655 0.097954 .
TskinThick  -0.002132   0.008245  -0.259 0.795976 .
Test        -0.000152   0.001099  -0.138 0.889997
BMI          0.075514   0.016796   4.496 6.92e-06 ***
DiabPediFunc 0.735885   0.351955   2.091 0.036542 *
Age          0.004790   0.011638   0.412 0.680616
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 683.88  on 533  degrees of freedom
Residual deviance: 514.94  on 525  degrees of freedom
AIC: 532.94

Number of Fisher Scoring iterations: 5

```

5.2.2 Interpretation

Factors to be checked –

- Variable Significance/Insignificance
- AIC(Akaike information criterion)
- Fisher Scoring Iterations

The variables – TskinThick, Test and Age seem to be statistically insignificant.

AIC(Akaike information criterion) - is an estimator of the relative quality of statistical models for a given set of data. It is another measure of goodness of fit that takes into account the ability of the model to fit the data. Thus, AIC provides a means for model selection. The lower the AIC, the better is the model.

Fisher scoring iterations uses an iterative approach (the Newton-Raphson algorithm by default) that looks for the best model. The algorithm stops when it doesn't perceive that moving again would yield much additional improvement. This line tells you how many iterations there were before the process stopped and output the results.

5.2.3 Multicollinearity Check

VIF(Variable Inflation Factor) – measures how much the variance of a regression coefficient is inflated due to multicollinearity in the model. It helps solve the problem of Multicollinearity(implies that the information that this variable provides about the response is redundant in the presence of the other variables).

```
> library(car)
> vif(logit_model1)
```

| NoPreg | PlaGluConc | DiastolicBP | TSkinThick | Test | BMI | DiabPediFunc | Age |
|----------|------------|-------------|------------|----------|----------|--------------|----------|
| 1.531278 | 1.179534 | 1.171907 | 1.535667 | 1.457756 | 1.192764 | 1.025955 | 1.614628 |

The result indicated that the variables aren't correlated to each other. This will help in building a model with good accuracy.

5.2.4 Improving the Model(AIC)

```
> logit_model2 <- glm(Class ~ NoPreg+PlaGluConc+BMI+DiabPediFunc+DiastolicBP, data = train_data,
+ family = binomial)
> summary(logit_model2)
```

Call:
glm(formula = Class ~ NoPreg + PlaGluConc + BMI + DiabPediFunc + DiastolicBP, family = binomial, data = train_data)

Deviance Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|---------|---------|--------|--------|
| -2.5922 | -0.7467 | -0.4477 | 0.7706 | 2.9109 |

Coefficients:

| | Estimate | Std. Error | z value | Pr(> z) |
|--------------|-----------|------------|---------|--------------|
| (Intercept) | -7.435422 | 0.782863 | -9.498 | < 2e-16 *** |
| NoPreg | 0.153437 | 0.031991 | 4.796 | 1.62e-06 *** |
| PlaGluConc | 0.032920 | 0.004023 | 8.182 | 2.78e-16 *** |
| BMI | 0.073267 | 0.015959 | 4.591 | 4.42e-06 *** |
| DiabPediFunc | 0.720630 | 0.348583 | 2.067 | 0.0387 * |
| DiastolicBP | -0.010090 | 0.006017 | -1.677 | 0.0936 . |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 683.88 on 533 degrees of freedom
Residual deviance: 515.31 on 528 degrees of freedom
AIC: 527.31

Number of Fisher Scoring iterations: 5

The AIC after removing the statistically insignificant variables is 527.31 which is better than the previous 532.94.

```

> step_model <- step(logit_model1)
Start: AIC=532.94
Class ~ NoPreg + PlaGluConc + DiastolicBP + TSkinThick + Test +
      BMI + DiabPediFunc + Age

  Df Deviance  AIC
- Test      1  514.96 530.96
- TSkinThick 1  515.00 531.00
- Age       1  515.11 531.11
<none>     1  514.94 532.94
- DiastolicBP 1  517.69 533.69
- DiabPediFunc 1  519.38 535.38
- NoPreg      1  529.65 545.65
- BMI         1  537.30 553.30
- PlaGluConc  1  584.52 600.52

Step: AIC=530.96
Class ~ NoPreg + PlaGluConc + DiastolicBP + TSkinThick + BMI +
      DiabPediFunc + Age

  Df Deviance  AIC
- TSkinThick 1  515.09 529.09
- Age       1  515.13 529.13
<none>     1  514.96 530.96
- DiastolicBP 1  517.70 531.70
- DiabPediFunc 1  519.38 533.38
- NoPreg      1  529.73 543.73
- BMI         1  537.43 551.43
- PlaGluConc  1  593.54 607.54

Step: AIC=529.09
Class ~ NoPreg + PlaGluConc + DiastolicBP + BMI + DiabPediFunc +
      Age

  Df Deviance  AIC
- Age       1  515.31 527.31
<none>     1  515.09 529.09
- DiastolicBP 1  518.08 530.08
- DiabPediFunc 1  519.40 531.40
- NoPreg      1  529.86 541.86
- BMI         1  538.70 550.70
- PlaGluConc  1  594.14 606.14

Step: AIC=527.31
Class ~ NoPreg + PlaGluConc + DiastolicBP + BMI + DiabPediFunc

  Df Deviance  AIC
<none>     1  515.31 527.31
- DiastolicBP 1  518.13 528.13
- DiabPediFunc 1  519.64 529.64
- BMI         1  538.71 548.71
- NoPreg      1  539.36 549.36
- PlaGluConc  1  598.86 608.86

```

We can use the step function in order to verify the selection of the best model with the lowest AIC score.

```

> anova(logit_model2, test = "Chisq")
Analysis of Deviance Table

Model: binomial, link: logit

Response: Class

Terms added sequentially (first to last)

      Df Deviance Resid. Df Resid. Dev Pr(>Chi)
NULL                                533      683.88
NoPreg      1   36.141      532      647.74 1.836e-09 ***
PlaGluConc  1  103.827      531      543.91 < 2.2e-16 ***
BMI         1   21.814      530      522.10 3.004e-06 ***
DiabPediFunc 1    3.973      529      518.13 0.04623 *
DiastolicBP  1    2.816      528      515.31 0.09330 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Likelihood Ratio – assesses the goodness of fit of two competing statistical models based on the ratio of their likelihoods. It compares the intercept-only model with the model with the predictor variables.

The p-value and Chisq value is significant, stating that the model built is significant enough for this model.


```
> lrtest(logit_model2)
Likelihood ratio test

Model 1: Class ~ NoPreg + PlaGluConc + BMI + DiabPediFunc + DiastolicBP
Model 2: Class ~ 1
#Df LogLik Df Chisq Pr(>Chisq)
1 6 -257.65
2 1 -341.94 -5 168.57 < 2.2e-16 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

This will show the robustness of the model. McFadden score, tells – how much of the variance is actually explained by this particular model over the intercept-only model.

```
> pr2(logit_model1)
      llh      llhNull      G2      McFadden      r2ML      r2CU
-257.4688426 -341.9401912 168.9426974 0.2470354 0.2712118 0.3755614
> 1-(257.4688426/341.9401912)
[1] 0.2470354
```

Coefficients of the variables and how much they are contributing to the overall model.
DiabPediFunc, NoPreg, BMI and PlaGluConc are contributing the most to the model.

```
> # Odds Ratio
> exp(coef(logit_model2))
(Intercept)      NoPreg      PlaGluConc      BMI DiabPediFunc      DiastolicBP
0.0005899801 1.1658338559 1.0334679672 1.0760176594 2.0557275065 0.9899607589
```

Probabilities for all the variables –

```
> # Probability
> exp(coef(logit_model2))/(1+exp(coef(logit_model2)))
(Intercept)      NoPreg      PlaGluConc      BMI DiabPediFunc      DiastolicBP
0.0005896322 0.5382840668 0.5082292831 0.5183085291 0.6727456889 0.4974775279
```

5.2.5 Tuning the Model(to improve Sensitivity)

```
> pred = predict(logit_model2, data=train_data, type="response")
> y_pred_num = ifelse(pred>0.5,1,0)
> y_pred = factor(y_pred_num, levels=c(0,1))
> y_actual = train_data$Class
> confusionMatrix(y_pred,train_data$Class,positive="1")
Confusion Matrix and Statistics

          Reference
Prediction 0      1
      0  312    80
      1   41   101

      Accuracy : 0.7734
      95% CI   : (0.7355, 0.8083)
 No Information Rate : 0.661
 P-Value [Acc > NIR] : 9.542e-09

      Kappa : 0.4663

McNemar's Test P-Value : 0.0005512

      Sensitivity : 0.5580
      Specificity : 0.8839
   Pos Pred Value : 0.7113
   Neg Pred Value : 0.7959
      Prevalence : 0.3390
   Detection Rate : 0.1891
 Detection Prevalence : 0.2659
   Balanced Accuracy : 0.7209

      'Positive' Class : 1
```

Sensitivity/True Positive Rate – These are cases in which we predicted yes (they have the disease), and they do have the disease.

In the above model the Accuracy is 77.34% but the Sensitivity is just 55.8%. In order to increase the value for sensitivity, a threshold of 0.35 was taken instead of the previous 0.50.

In the new model though the accuracy has decreased by a small amount, the Sensitivity has drastically improved by 14.9%. We will therefore consider the revised threshold to determine the model True Positive Rate.

```

> # Calibrating threshold levels to increase sensitivity
> pred = predict(logit_model2, data=train_data, type="response")
> y_pred_num = ifelse(pred>0.35,1,0)
> y_pred = factor(y_pred_num, levels=c(0,1))
> y_actual = train_data$Class
> confusionMatrix(y_pred,y_actual,positive="1")
Confusion Matrix and Statistics

      Reference
Prediction 0    1
0      277   53
1       76  128

      Accuracy : 0.7584
      95% CI   : (0.7198, 0.7942)
    No Information Rate : 0.661
    P-Value [Acc > NIR] : 6.583e-07

      Kappa : 0.4771

McNemar's Test P-Value : 0.05275

      Sensitivity : 0.7072
      Specificity : 0.7847
    Pos Pred Value : 0.6275
    Neg Pred Value : 0.8394
      Prevalence : 0.3390
    Detection Rate : 0.2397
Detection Prevalence : 0.3820
    Balanced Accuracy : 0.7459

'Positive' Class : 1

```

5.3 K-NN(K-Nearest Neighbour)

5.3.1 Normalizing continuous variables

Distance metric is highly influenced by the scale of the variable. Hence, it is important to standardize variables before utilizing them in model building. We will use min-max standardization method to bring all variables in same scale.

5.3.2 Model Building

```

> scale = preProcess(train_data, method = "range")
> train.norm.data = predict(scale, train_data)
> test.norm.data = predict(scale, test_data)
> knn_fit = train(Class ~., data = train.norm.data, method = "knn",
+               trControl = trainControl(method = "cv", number = 3),
+               tuneLength = 10)
> knn_fit
k-Nearest Neighbors

534 samples
  8 predictor
  2 classes: '0', '1'

No pre-processing
Resampling: Cross-Validated (3 fold)
Summary of sample sizes: 356, 356, 356
Resampling results across tuning parameters:

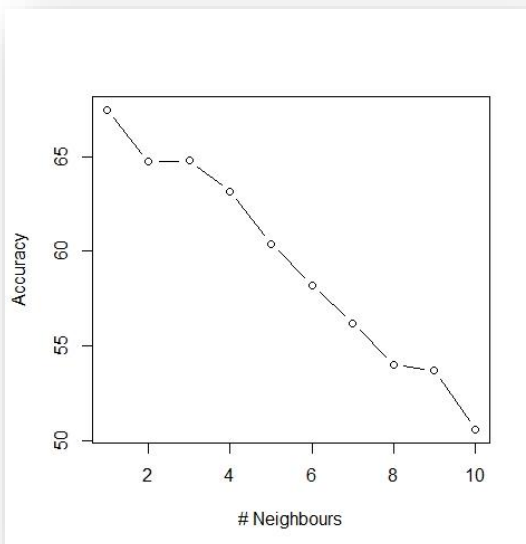
 k  Accuracy  Kappa
  5  0.7209738 0.3309169
  7  0.7284644 0.3585581
  9  0.7322097 0.3674254
 11  0.7322097 0.3584005
 13  0.7209738 0.3255866
 15  0.7265918 0.3244056
 17  0.7284644 0.3320548
 19  0.7322097 0.3392210
 21  0.7340824 0.3468878
 23  0.7322097 0.3355015

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was k = 21.

```

We use the “train” function to get the best k value from the model. It gives the best k value that was selected for this particular model.

This indicates that if we select 21 neighbours to make a classification decision, then this model will fit.



5.4 Naïve Bayes

5.4.1 Model Building

```
> NB
```

Naive Bayes Classifier for Discrete Predictors

Call:

```
naiveBayes.default(x = train.norm.data[-c(9)], y = train.norm.data$Class)
```

A-priori probabilities:

```
train.norm.data$Class
```

```
      0      1  
0.6610487 0.3389513
```

Conditional probabilities:

```
      NoPreg  
train.norm.data$Class      [,1]      [,2]  
0 0.1881353 0.1816886  
1 0.3012675 0.2251628
```

```
      PlaGluConc  
train.norm.data$Class      [,1]      [,2]  
0 0.5567363 0.1371966  
1 0.7115854 0.1578604
```

```
      DiastolicBP  
train.norm.data$Class      [,1]      [,2]  
0 0.5567965 0.1505474  
1 0.5858618 0.1682860
```

```
      TSkinThick  
train.norm.data$Class      [,1]      [,2]  
0 0.3104006 0.2376306  
1 0.3359642 0.2711273
```

```
      Test  
train.norm.data$Class      [,1]      [,2]  
0 0.08789104 0.1299195  
1 0.13477990 0.1839854
```

```
      BMI  
train.norm.data$Class      [,1]      [,2]  
0 0.4540388 0.1161715  
1 0.5200863 0.1149361
```

```
      DiabPediFunc  
train.norm.data$Class      [,1]      [,2]  
0 0.1541309 0.1343762  
1 0.1955829 0.1527359
```

```
      Age  
train.norm.data$Class      [,1]      [,2]  
0 0.1709160 0.1948376  
1 0.2642726 0.1762445
```

Probabilities identified across all the variables

6. Model Comparison

6.1 Logistic Regression

6.1.1 Train Dataset

```
> # Calibrating threshold levels to increase sensitivity
> pred = predict(logit_model2, data=train_data, type="response")
> y_pred_num = ifelse(pred>0.35,1,0)
> y_pred = factor(y_pred_num, levels=c(0,1))
> y_actual = train_data$Class
> confusionMatrix(y_pred,y_actual,positive="1")
Confusion Matrix and Statistics

      Reference
Prediction 0    1
0      277    53
1       76   128

      Accuracy : 0.7584
      95% CI   : (0.7198, 0.7942)
    No Information Rate : 0.661
    P-Value [Acc > NIR] : 6.583e-07

      Kappa : 0.4771

McNemar's Test P-Value : 0.05275

      Sensitivity : 0.7072
      Specificity : 0.7847
    Pos Pred Value : 0.6275
    Neg Pred Value : 0.8394
      Prevalence : 0.3390
    Detection Rate : 0.2397
    Detection Prevalence : 0.3820
    Balanced Accuracy : 0.7459

      'Positive' Class : 1
```

6.1.2 Test Data

```
> # Performance metrics (test sample)
> pred = predict(logit_model2, newdata=test_data, type="response")
> y_pred_num = ifelse(pred>0.35,1,0)
> y_pred = factor(y_pred_num, levels=c(0,1))
> y_actual = test_data$class
> confusionMatrix(y_pred,y_actual,positive="1")
Confusion Matrix and Statistics
```

| | Reference | |
|------------|-----------|----|
| Prediction | 0 | 1 |
| 0 | 122 | 23 |
| 1 | 25 | 64 |

```

          Accuracy : 0.7949
          95% CI : (0.7374, 0.8447)
    No Information Rate : 0.6282
    P-Value [Acc > NIR] : 2.749e-08

          Kappa : 0.5629

McNemar's Test P-Value : 0.8852

    Sensitivity : 0.7356
    Specificity : 0.8299
    Pos Pred Value : 0.7191
    Neg Pred Value : 0.8414
    Prevalence : 0.3718
    Detection Rate : 0.2735
    Detection Prevalence : 0.3803
    Balanced Accuracy : 0.7828

'Positive' Class : 1

```

6.2 K Nearest Neighbour

6.2.1 Train Dataset

```
> # Performance metrics (train sample)
> pred = predict(knn_fit, data = train.norm.data[-9], type = "raw")
> confusionMatrix(pred,train.norm.data$class,positive="1")
Confusion Matrix and Statistics
```

| | Reference | |
|------------|-----------|----|
| Prediction | 0 | 1 |
| 0 | 322 | 96 |
| 1 | 31 | 85 |

```

          Accuracy : 0.7622
          95% CI : (0.7237, 0.7977)
    No Information Rate : 0.661
    P-Value [Acc > NIR] : 2.427e-07

          Kappa : 0.4184

McNemar's Test P-Value : 1.354e-08

    Sensitivity : 0.4696
    Specificity : 0.9122
    Pos Pred Value : 0.7328
    Neg Pred Value : 0.7703
    Prevalence : 0.3390
    Detection Rate : 0.1592
    Detection Prevalence : 0.2172
    Balanced Accuracy : 0.6909

'Positive' Class : 1

```

6.2.2 Test Data

```
> # Performance metrics (test sample)
> pred = predict(knn_fit, newdata = test.norm.data[-9], type = "raw")
> confusionMatrix(pred, test.norm.data$Class, positive = "1")
Confusion Matrix and Statistics

      Reference
Prediction 0  1
0      134  49
1       13  38

      Accuracy : 0.735
      95% CI   : (0.6736, 0.7904)
    No Information Rate : 0.6282
    P-Value [Acc > NIR] : 0.0003513

      Kappa : 0.3805

  Mcnemar's Test P-Value : 8.789e-06

    Sensitivity : 0.4368
    Specificity : 0.9116
   Pos Pred Value : 0.7451
   Neg Pred Value : 0.7322
    Prevalence : 0.3718
    Detection Rate : 0.1624
    Detection Prevalence : 0.2179
    Balanced Accuracy : 0.6742

 'Positive' Class : 1
```

6.3 Naïve Bayes

6.3.1 Train Dataset

```
> # Performance metrics (train sample)
> pred_NB_train = predict(NB, newdata = train.norm.data[-9])
> confusionMatrix(pred_NB_train, train.norm.data$Class, positive = "1")
Confusion Matrix and Statistics

      Reference
Prediction 0  1
0      300  75
1       53 106

      Accuracy : 0.7603
      95% CI   : (0.7218, 0.7959)
    No Information Rate : 0.661
    P-Value [Acc > NIR] : 4.017e-07

      Kappa : 0.4488

  Mcnemar's Test P-Value : 0.06343

    Sensitivity : 0.5856
    Specificity : 0.8499
   Pos Pred Value : 0.6667
   Neg Pred Value : 0.8000
    Prevalence : 0.3390
    Detection Rate : 0.1985
    Detection Prevalence : 0.2978
    Balanced Accuracy : 0.7177

 'Positive' Class : 1
```


6.3.2 Test Data

```
> # Performance metrics (test sample)
> pred_NB_test = predict(NB, newdata = test.norm.data[-9])
> confusionMatrix(pred_NB_test, test.norm.data$Class, positive="1")
Confusion Matrix and Statistics

              Reference
Prediction    0      1
0      127    35
1       20    52

              Accuracy : 0.765
              95% CI   : (0.7053, 0.8178)
              No Information Rate : 0.6282
              P-Value [Acc > NIR] : 5.444e-06

              Kappa : 0.4785

McNemar's Test P-Value : 0.05906

              Sensitivity : 0.5977
              Specificity : 0.8639
              Pos Pred Value : 0.7222
              Neg Pred Value : 0.7840
              Prevalence : 0.3718
              Detection Rate : 0.2222
              Detection Prevalence : 0.3077
              Balanced Accuracy : 0.7308

              'Positive' Class : 1
```

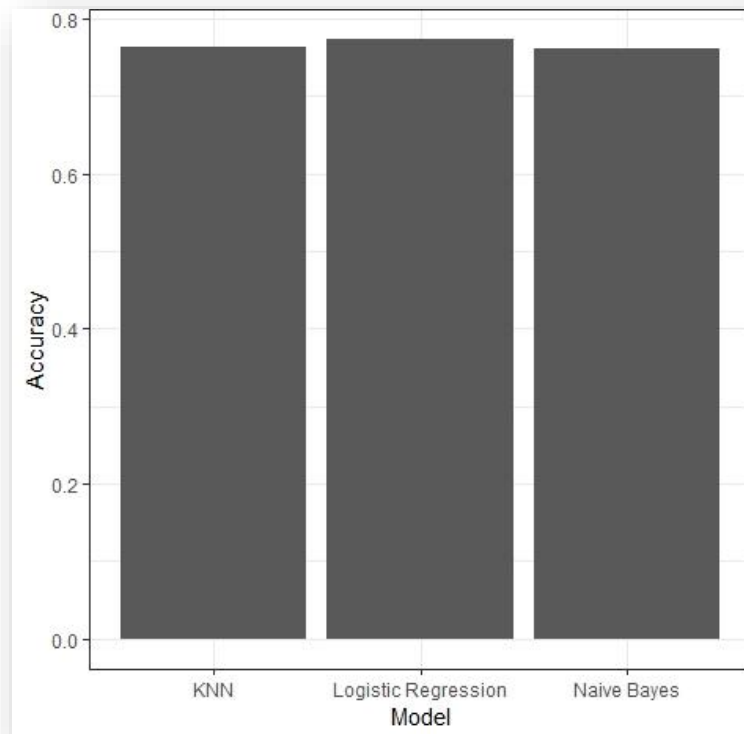
For Logistic Regression, KNN and Naïve Bayes, the model when validated on the test/validation dataset gives a result similar to that of the train dataset. We can therefore say that the model is not an over-fit model.

7. Cross Validation & Insights

7.1 Cross Validation

In this problem, decision makers will be keen to identify the positives very accurately. Therefore, we will not just evaluate the model based on the Accuracy but also use Sensitivity as a measuring factor to compare the models.

| | Accuracy | | | Sensitivity | | | Specificity | | |
|-------|---------------------|-----|-------------|---------------------|-----|-------------|---------------------|-----|-------------|
| | Logistic Regression | KNN | Naïve Bayes | Logistic Regression | KNN | Naïve Bayes | Logistic Regression | KNN | Naïve Bayes |
| Train | 77% | 76% | 76% | 71% | 47% | 59% | 79% | 91% | 85% |
| Test | 79% | 74% | 77% | 75% | 44% | 60% | 82% | 91% | 86% |



From the above results it is clear that although there is not much difference in all the three algorithms, **Logistic Regression** gives a slightly better results at predicting the true positives as compared to the others.

The Diabetes Database was analysed and explored in detail. The patterns identified using Data Exploration methods were validated using the modelling techniques employed. Classification models such as Logistic Regression, K-Nearest Neighbour and Naïve Bayes were built and evaluated to identify best model to predict the occurrence of Diabetes in women. From the cross-validated performance measure of sensitivity, the Logistic

Regression model was concluded as the best performing model while Naïve Bayes seems to have performed the worst.

7.2 Insights

It is clear from the analysis that a marginal increase in most of the factors can increase the probability of a woman being diagnosed with diabetes. It is therefore, important to focus on correctly classifying the results over the misclassifications.

Therefore, the occurrence of the disease is significantly affected by the following factors -

- **NoPreg** – nbm,
- **PlaGluConc** –
- **DiastolicBP** –
- **TSkinThick** –
- **Test** –
- **BMI** –
- **DiabPediFunc** –
- **Age** –

8. Source of Data

- Great Learning Mentored Learning Session and Recorded Sessions
- Google
- R Blogger
- stats.stackexchange.com
- uc-r.github.io
- www.kaggle.com
- www.rpubs.com
- www.datacamp.com